

# HW 1

Nicole Sullivan

```
knitr::opts_chunk$set(echo = T, message = F, warning = F)
```

```
# DO THESE FIRST BEFORE LOADING RETICULATE PKG
# create a new environment
#reticulate::conda_create("hw1")
# install numpy
#reticulate::conda_install("hw1", "numpy")
# reticulate::conda_install("hw1", "scipy")

reticulate::use_condaenv("hw1")
```

```
library(tidyverse)
library(reticulate) # to use Python

# Plotting aesthetics
loaded_font <- "Didact Gothic"
text_color <- "white"

# Hex codes
hex_purple <- "#5E72E4" #primary
hex_blue_lt <- "#5DC2F8" #info
hex_green <- "#66BB6A" #success
hex_pink <- "#F44336" #danger
hex_orange <- "#FF9800" #warning
hex_blue_dk <- "#17724D" #default
hex_grey <- "#545454"
hex_blue_deep <- "#00151C"

proj_theme <- theme(plot.background = element_rect(fill = "#1c2030", color = "transparent"),
  plot.margin = margin(t = "1.5", r = "1.5", b = "1.5", l = "1.5", unit = "cm"),
  panel.background = element_rect(fill = "#1c2030"),
  panel.grid.major = element_blank(),
  panel.grid.minor = element_blank(),
  plot.title = element_text(family = loaded_font, color = text_color, hjust = 0.5, face = "bold", size = 14),
  plot.caption = element_text(family = loaded_font, color = text_color, size = 8),
  axis.title = element_text(family = loaded_font, size = 15, color = text_color),
  axis.text = element_text(family = loaded_font, color = text_color, size = 10),
  strip.background = element_rect(fill = "transparent"),
  strip.text = element_text(color = "white", family = loaded_font, face = "bold", size = 10),
  legend.background = element_rect(fill = "transparent"),
  legend.title = element_text(family = loaded_font, color = text_color),
  legend.text = element_text(family = loaded_font, color = text_color),
  legend.position = "bottom",
  legend.key = element_rect(fill = NA))
```

## 1

Find the Maximum Likelihood Estimation (MLE) of  $\theta$  in the following probabilistic density functions. In each case, consider a random sample of size  $n$ . Show your calculation.

### 1a

$$f(x|\theta) = xe^{-x\theta}, x \geq 0$$

$$L(\theta|x) = \prod p(x_i|\theta)$$

Take the log:

$$= \sum \log p(x_i|\theta)$$

Plug in the likelihood:

$$= \sum \log \left( \frac{x}{\theta} e^{-\frac{x^2}{\theta}} \right)$$

Distribute the logs:

$$= \sum \left( \log \frac{x}{\theta^2} - \frac{x^2}{2\theta^2} \right)$$

Distribute the summation:

$$= \sum \log \frac{x}{\theta^2} - \frac{\sum x^2}{2\theta^2}$$

Use log rules to rewrite:

$$= \sum \log x - \sum \log \theta^2 - \frac{\sum x^2}{2\theta^2}$$

Now, we take the derivative w.r.t.  $\theta$ :

$$\begin{aligned} \frac{\partial}{\partial \theta} &= 0 - \frac{2N}{\theta} - \frac{\sum x^2}{\theta^3} \\ &= -\frac{2N}{\theta} + \frac{\sum x^2}{\theta^3} \end{aligned}$$

Set equal to zero to maximize:

$$\begin{aligned} 0 &= -\frac{2N}{\theta} + \frac{\sum x^2}{\theta^3} \\ \frac{2N}{\theta} &= \frac{\sum x^2}{\theta^3} \\ 2N\theta^2 &= \sum x^2 \\ \hat{\theta} &= \sqrt{\frac{\sum x^2}{2N}} \end{aligned}$$

### 1b

$$f(x|\theta, \alpha, \beta) = a\theta^{-\alpha}\beta x^{\beta}e^{-\frac{1}{\theta}x^{\beta}}$$

Plug in likelihood and take logs:

$$\begin{aligned} \sum \log(a\theta^{-\alpha\beta}) + \sum \log x^{\beta} - \sum \left( \frac{x^{\beta}}{\theta} \right) \\ \sum \log(a\theta^{-\alpha\beta}) + \sum \log x^{\beta} - \frac{\sum x^{\beta}}{\theta^{\beta}} \end{aligned}$$

Next, we take the derivative w.r.t.  $\theta$  ( $\alpha$  and  $\beta$  are hyperparameters):

$$\begin{aligned} \frac{\partial}{\partial \theta} &= \frac{N\alpha}{\theta^{-\alpha\beta-1}} \cdot a\theta + 0 + \beta \sum x^{\beta} \theta^{-\beta-1} \\ &= \frac{-N\alpha^2\beta}{\theta^{-\alpha\beta-1}} + \beta \sum x^{\beta} \theta^{-\beta-1} \end{aligned}$$

Now, set equal to zero.

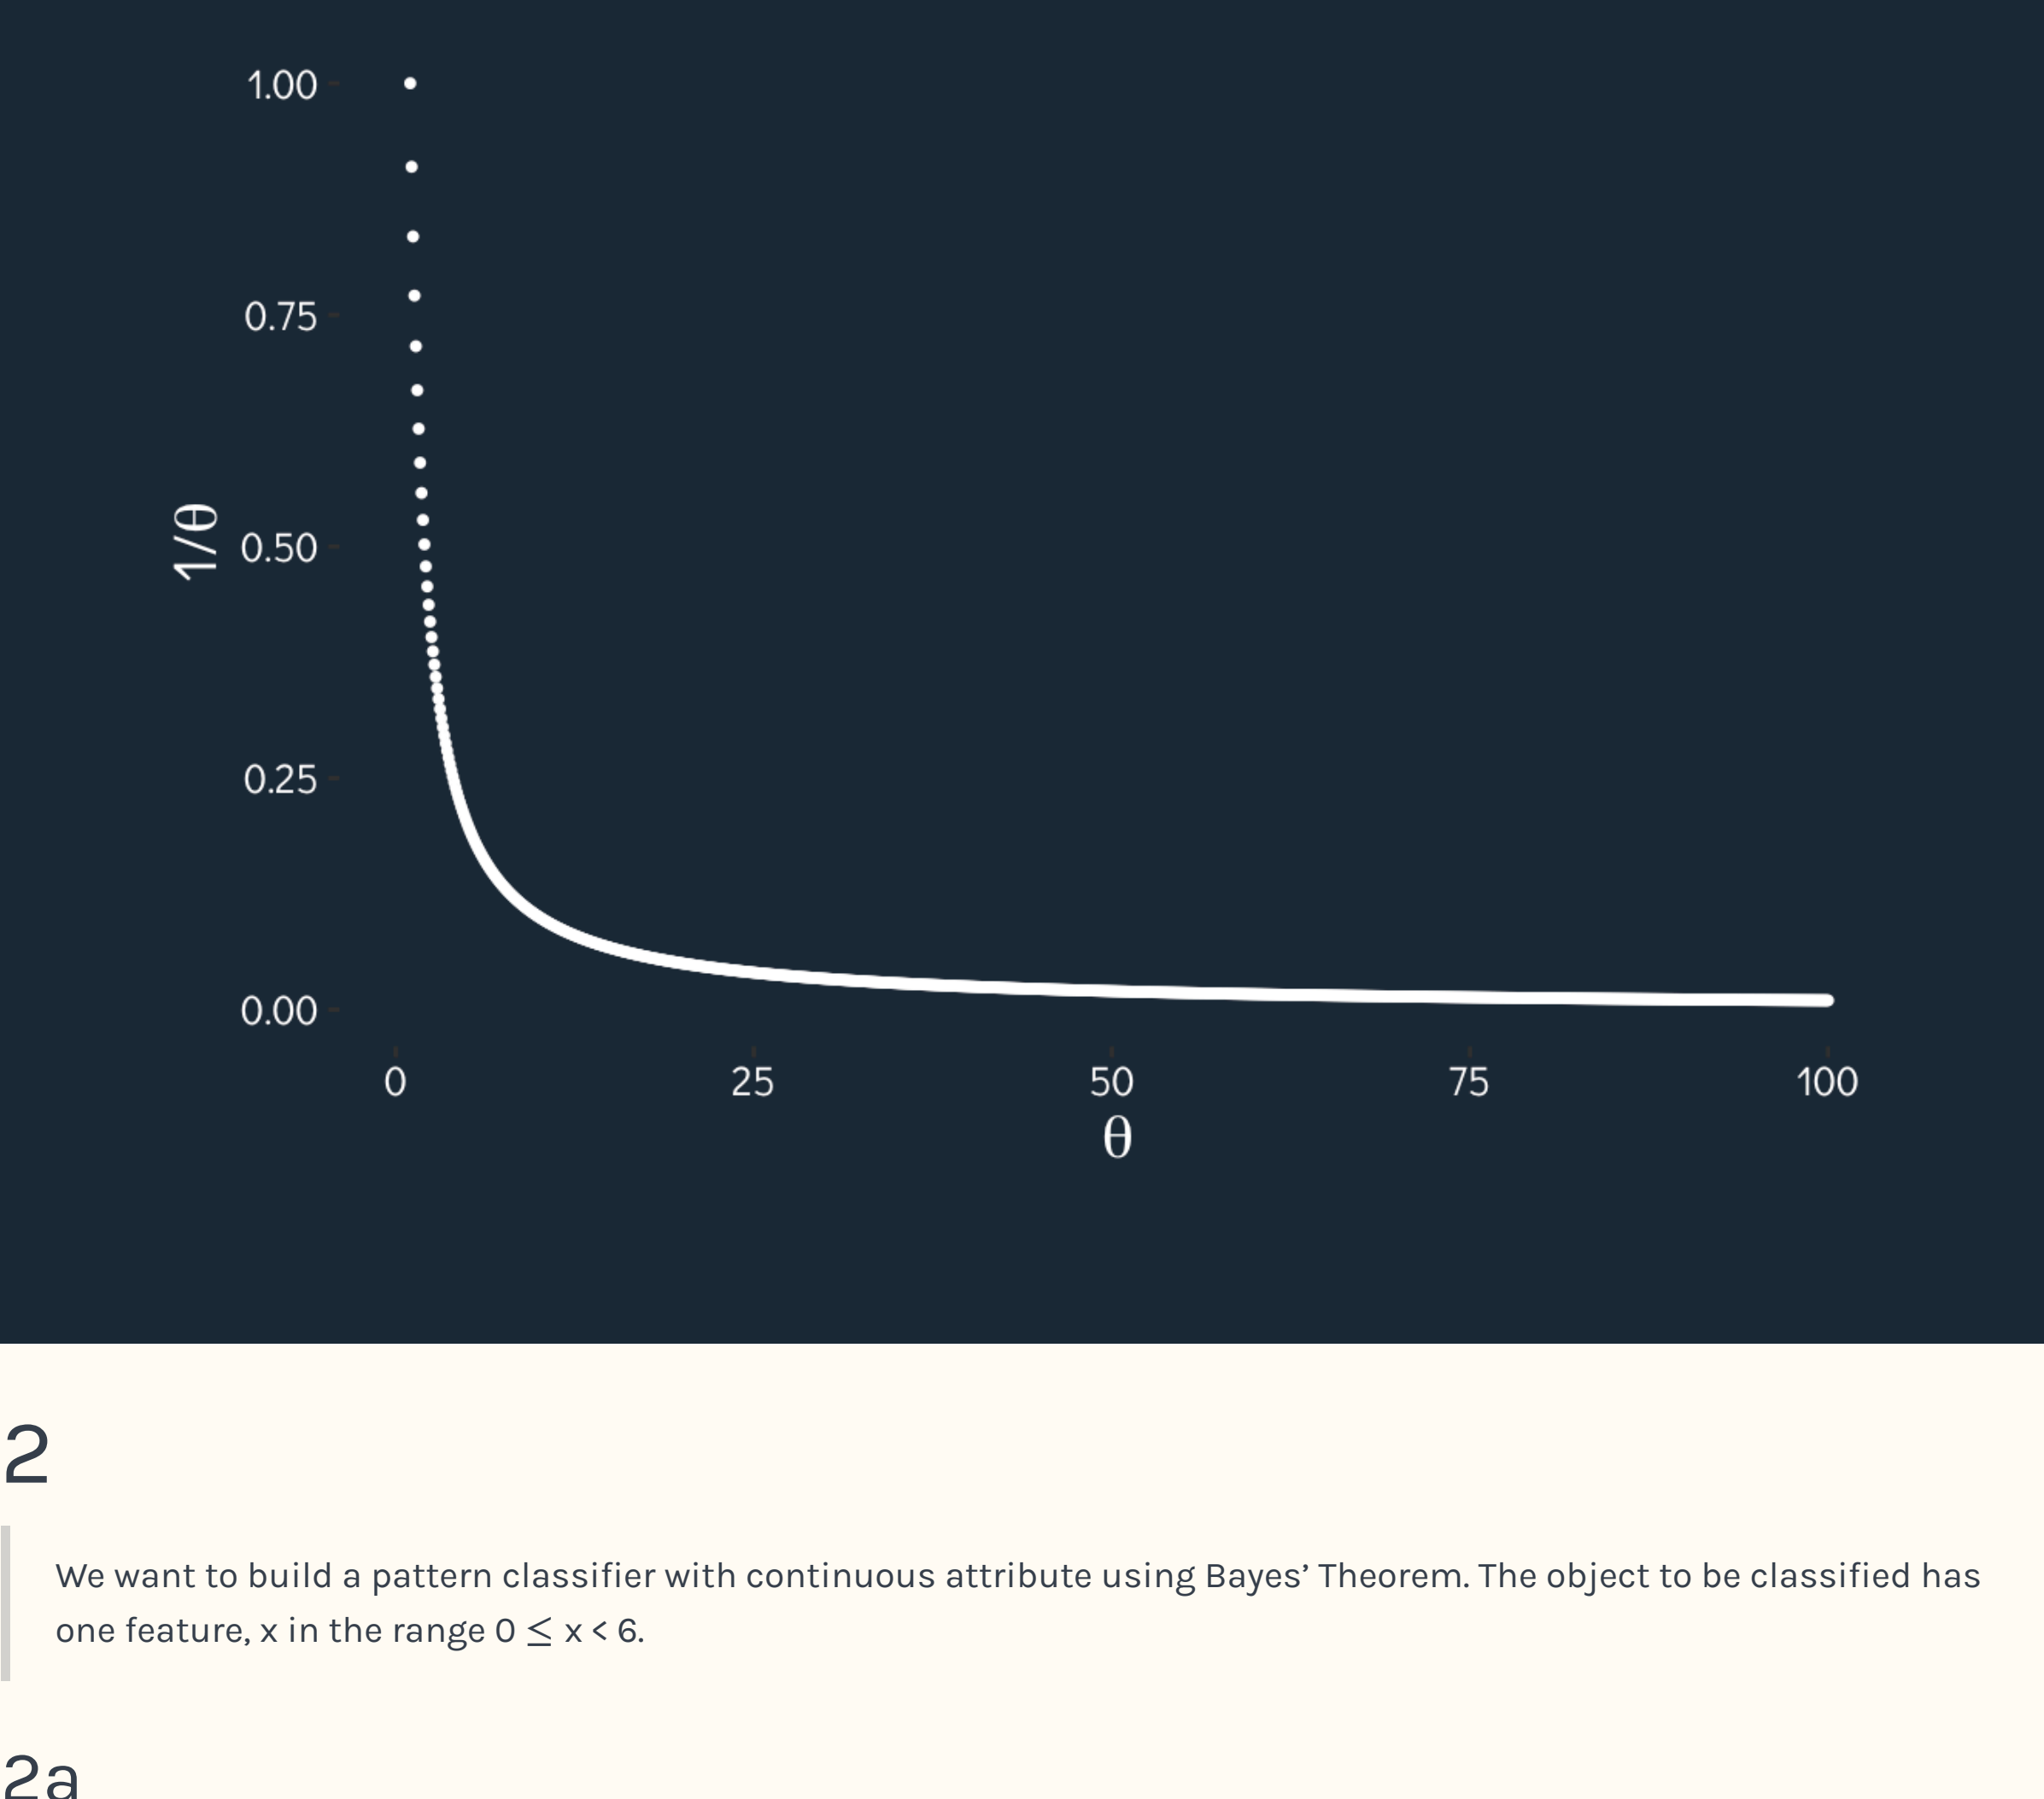
$$\begin{aligned} 0 &= \frac{-N\alpha^2\beta}{\theta^{-\alpha\beta-1}} + \beta \sum x^{\beta} \theta^{-\beta-1} \\ \frac{N\alpha^2\beta}{\theta^{-\alpha\beta-1}} &= \beta \sum x^{\beta} \theta^{-\beta-1} \\ \frac{N\alpha^2\beta}{\theta^{-\alpha\beta-1}} &= \beta \frac{\sum x^{\beta}}{\theta^{\beta+1}} \\ \frac{N\alpha^2\beta\theta^{\beta+1}}{\theta^{-\alpha\beta-1}} &= \beta \sum x^{\beta} \\ \theta^{(\beta+1)-(-\alpha\beta-1)} &= \frac{\sum x^{\beta}}{N\alpha^2} \\ \theta^{\beta-\alpha\beta+2} &= \frac{\sum x^{\beta}}{N\alpha^2} \\ \hat{\theta} &= \left( \frac{\sum x^{\beta}}{N\alpha^2} \right)^{\frac{1}{\beta-\alpha\beta+2}} \end{aligned}$$

### 1c

$$f(x|\theta) = \frac{1}{\theta}, 0 \leq x \leq \theta, \theta > 0$$

This function has no minimum or maximum, as it is monotonically decreasing.

To illustrate, we can also draw (or plot) the function:



## 2

We want to build a pattern classifier with continuous attribute using Bayes' Theorem. The object to be classified has one feature,  $x$  in the range  $0 \leq x \leq 6$ .

### 2a

Assuming equal priors,  $P(C1) = P(C2) = 0.5$ , classify an object with the attribute value  $x = 2.5$ .

To classify, I'll determine the value of the discriminant for each class for this value of  $x$ . Then I'll classify the observation as the class with the highest discriminant.

**Class 1:**

$$P(x|C_1) \times P(C_1) = \frac{1}{6} \times \frac{1}{2} = \frac{1}{12} = 0.08\bar{3}$$

**Class 2:**

$$\begin{aligned} P(x|C_2) \times P(C_2) &= \frac{1}{4} (2.5 - 1) \times 0.5 \\ &= \frac{1}{4} (1.5) \times \frac{1}{2} = \frac{1}{4} \times \frac{3}{2} \times \frac{1}{2} = \frac{3}{16} = 0.1875 \end{aligned}$$

$P(C_2|x) > P(C_1|x)$   $\therefore$  I'll classify as  $C_2$ .

### 2b

Assuming unequal priors,  $P(C1) = 0.7, P(C2) = 0.3$ , classify an object with the attribute value  $x = 4$ .

**Class 1:**

$$p(x|C_1) \times p(C_1) = \frac{1}{6} \times 0.7 = 0.11\bar{6}$$

**Class 2:**

$$\begin{aligned} p(x|C_2) \times p(C_2) &= \frac{1}{4} (5 - 4) \times 0.3 \\ &= \frac{1}{4} \times 0.3 = 0.075 \end{aligned}$$

$p(C_1|x) > p(C_2|x)$   $\therefore$  I'll classify as  $C_1$ .

### 2c

Consider a decision function  $\phi(x)$  of the form  $\phi(x) = (x - 3) - \alpha$  with one free parameter  $\alpha$  in the range  $0 \leq \alpha \leq 2$ . You classify a given input  $x$  as class 2 if and only if  $\phi(x) < 0$ , or equivalently  $3 - \alpha < x < 3 + \alpha$ , otherwise you choose  $x$  as class 1. Assume equal priors,  $P(C1) = P(C2) = 0.5$ , what is the optimal decision boundary - that is, what is the value of  $\alpha$  which minimizes the probability of misclassification? What is the resulting probability of misclassification with this optimal value for  $\alpha$ ? (Hint: take advantage of the symmetry around  $x = 3$ )

We basically treat this as a minimization problem: the optimal  $\alpha$  will minimize **total misclassification**. To get total misclassification, we need to sum:

- $P(C_2|C_1)$ : the probability of classifying as  $C_2$  when the observation is actually  $C_1$ , and
- $P(C_1|C_2)$ : the probability of classifying as  $C_1$  when the observation is really  $C_2$ .

Or, more succinctly put:

$$\text{Total misclassification rate} = P(C_2|C_1) + P(C_1|C_2)$$

where:

- $P(C_2|C_1) = \int_{3-\alpha}^{3+\alpha} f_1(x) dx$
- $P(C_1|C_2) = \int_{3-\alpha}^{3+\alpha} f_2(x) dx$

or, in plain English, the misclassification rate for, say,  $C_1$  is the total area under the overlap of the  $C_2$  and  $C_1$  in the  $C_2$  classification range.

That means that:

- $P(C_2|C_1) = \int_{3-\alpha}^{3+\alpha} \frac{1}{4} dx$
- $P(C_1|C_2) = \int_{3-\alpha}^{3+\alpha} \frac{1}{4} (x - 1) dx + \int_{3+\alpha}^6 \frac{1}{4} (5 - x) dx$

and:

$$TMR = \int_{3-\alpha}^{3+\alpha} \frac{1}{6} dx + \int_{3-\alpha}^3 \frac{1}{4} (x - 1) dx + \int_{3+\alpha}^6 \frac{1}{4} (5 - x) dx$$

Integrating w.r.t  $x$  we get:

$$TMR = \left[ \frac{1}{6} x \right]_{3-\alpha}^{3+\alpha} + \left[ \frac{1}{4} x^2 + \frac{1}{4} x \right]_{3-\alpha}^{3+\alpha} + \left[ \frac{1}{4} x^2 - \frac{5}{4} x \right]_{3+\alpha}^6$$

and then evaluating the integral gives:

$$TMR = \frac{1}{8} \alpha^2 + \frac{3}{8} + \frac{1}{24} (-4\alpha + 3(\alpha - 3)(\alpha - 1))$$

To minimize misclassification, we take the derivative w.r.t.  $\alpha$ :

$$\frac{d}{d\alpha} TMR = \frac{\alpha}{2} - \frac{2}{3}$$

and then set equal to zero:

$$0 = \frac{\alpha}{2} - \frac{2}{3}$$

and solve for  $\alpha$ :

$$\alpha = \frac{4}{3}$$

**Therefore, we know that the value of  $\alpha$  that minimizes misclassification is  $\frac{4}{3}$ .**

To determine the rate of misclassification that will result from this decision boundary, we just need to plug  $\frac{4}{3}$  back into our original equation for TMR:

$$\frac{1}{8} \left( \frac{4}{3} \right)^2 + \frac{3}{8} + \frac{1}{24} \left( -4 \cdot \frac{4}{3} + 3 \cdot \frac{4}{3} - 3 \cdot \frac{4}{3} \left( \frac{4}{3} - 1 \right) \right)$$

which can easily be evaluated with code:

```
alpha <- 4/3

(1/8) * alpha^2 + (3/8) + ( (1/24) * (-4 * alpha + 3 * (alpha - 3) * (alpha - 1)))

## [1] 0.3855556
```

or, there's about a **31% total misclassification rate** when using  $\alpha = \frac{4}{3}$ .

Just for fun, we can also confirm these results by integrating/minimizing programmatically.

### Method 1: Minimize the overlap between the pdf's

The first method just uses intuition: to minimize misclassification, we just need to minimize the overlap of the pdf's of the two classes, or essentially, maximize the difference between the two:

$$\text{argmax} \int P(x|C_1) - \int P(x|C_2))$$



Our misclassification rate is minimized at  $\alpha = \frac{4}{3}$ , or when  $\frac{1}{3} < x < \frac{5}{3}$ .

```
c2_misclass = quad(c2, a = 1, b = 5/3)[1] + quad(c2, a = 13/3, b = 5)[1]
pc2_misclass = c2_misclass * 0.5 # 0.5 prior prob
pc1_misclass = quad(c1, a = 1/3, b = 13/3)[1] * 0.5 # 0.5 prior prob
```

### Method 2: integrate to find total misclassification for each $\alpha$ & minimize

```
def p_misclass(alpha):
  lwr = 3 - alpha
  upr = 3 + alpha

  c1_misclass = quad(c1, a = lwr, b = upr)[1]
  c2_misclass = quad(c2, a = lwr, b = upr)[1] + quad(c2, a = upr, b = 6)[1]

  return c1_misclass + c2_misclass
```

```
misclass_result = []
alpha_vals = np.linspace(0, 2, 100)

for i in alpha_vals:
  misclass_result.append(p_misclass(i))
```

```
min_misclass = alpha_vals[np.array(misclass_result).argmin()]
# min misclass
# np.array(misclass_result).argmin() * 0.5 # x prior
```

```
ggplot() +
  aes(as.numeric(py$alpha_vals), as.numeric(py$misclass_result) * 0.5) +
  geom_point(size = 0.7, color = hex_purple) +
  geom_vline(xintercept = py$min_misclass, color = "white", size = 1.2) +
  labs(title = "Method 2: Minimize misclassification",
    x = expression(alpha),
    y = "Misclassifications") +
  proj_theme
```



We again find the optimal  $\alpha$  is  $\frac{4}{3}$ !